


With outer AdamW, adaptation effects are small

Top: extra target learning. Bottom: retention gain. Positive favors adaptation; numbers are in units of 10^{-5} NLL.



All 48 adaptation comparisons. Maximum retention gain = $3.64e-05$ NLL; none reaches the predeclared 0.001 screening threshold. Outlines mark 5 cells with both marginal 95% paired-document intervals below zero for adapted minus zero; effects remain small. All 16 own-zero comparisons to ordinary AdamW have identical document scores. Two-step reset-Adam inner adaptation, every update. 512 documents/domain, 5,000 paired bootstrap draws; one training seed and examined holdouts. No joint-95% or multiplicity guarantee.