


CPT retention gains must survive tuned ordinary controls



All 96 accepted terminals, one training seed. Better is right and down. All models still forget relative to base. Top four axes are linear within ± 0.2 , logarithmic beyond; bottom windows are linear and state their coverage. Math focus candidates are dominated by ordinary AdamW. Code shows a sampled tradeoff, not a matched-learning advantage. Regimes change both unique-data coverage and repetition. Compute and tuning budgets differ; no seed-uncertainty claim.