


Standard feedback: small, setting-dependent differences from own zero

Negative favors adaptation. Each vertical scale is fitted to its displayed effects; units are identical.



All 16 fixed-checkpoint 1× comparisons; both outcomes shown. Error bars: 95% paired-document bootstrap intervals (5,000 draws). Intervals are conditional on one trained trajectory; they do not measure training-seed variability or adjust for selection/multiple comparisons. No terminal candidate passes the frozen practical / tuned-ordinary-frontier screen. A small own-zero benefit is not tuned-baseline superiority. Four 10× training collapses remain fully reported in the companion curve figure. These results concern this implementation and finite grid.