



Does finer ordinary CPT tuning explain the retention advantage?


A SmolLM2-135M: ordinary training curves




B Effects after baseline tuning



C SmolLM2-360M: ordinary training curves



D Effects after baseline tuning



—●— New ordinary LR 7.5e-06 - - - Previous ordinary LR 1e-5 ★ Fixed Lookahead K64
—●— New ordinary LR 8.75e-06 ● Other measured ordinary controls ▲ Strongest compatible ordinary
—●— New ordinary LR 1.25e-05

Six new ordinary trajectories; all 270 planned checkpoints scored. Lines join measured points only.
 Comparator: lowest FineWeb NLL among ordinary points with math NLL $\leq$ candidate +0.001.
 One training seed; reused 512 math and 512 FineWeb documents. Pointwise CIs do not adjust for comparator selection.
 Negative FineWeb change means improvement over base. Lookahead-style interpolation is reported separately from T³.