


Terminal comparisons of feedback operators



All 16 prescribed terminal operator comparisons. Bars: 95% paired-document bootstrap intervals; no new statistical calculations in this figure. Negative differences favor the left method. The left panel uses a symlog axis to show both effect scales; the right panel is linear. Intervals condition on one training seed and previously examined documents; no correction for selection or multiple comparisons. Adapted/frozen teachers also differ in anchoring history and support exposure. These contrasts alone do not establish superiority to tuned ordinary CPT.