


Smaller adaptation: learning and forgetting


Four new trajectories; 12 exact accepted controls

- Ordinary training
- Normalized gradient, 0.1×
- ×·· Zero adaptation
- ▲— Normalized gradient, 1×


Code · outer LR 3e-06




Math · outer LR 3e-06



Code · outer LR 1e-05



Math · outer LR 1e-05



All eight fixed checkpoints and the shared public base.

Nats per target token; reused controls are not new training seeds.